

figure_lstm

October 15, 2024

```
[26]: #
import numpy as np #
import pandas as pd # CSV /
import matplotlib.pyplot as plt #
import os #
print(os.listdir("./input")) # ./input

['test.csv', 'train.csv', 'sample_submission.csv']
```

```
[27]: # PyTorch
import torch # PyTorch
import torch.nn as nn #
from torch.autograd import Variable #
from sklearn.model_selection import train_test_split #
from torch.utils.data import DataLoader, TensorDataset #
```

```
[28]: #
# CSV float32
train = pd.read_csv("./input/train.csv", dtype=np.float32)

# 0-9
targets_numpy = train.label.values # 0-9
features_numpy = train.loc[:, train.columns != "label"].values / 255 #
# 0-1

# 80% 20%
features_train, features_test, targets_train, targets_test = train_test_split(
    features_numpy, targets_numpy, test_size=0.2, random_state=42)

# PyTorch Tensor
featuresTrain = torch.from_numpy(features_train) #
targetsTrain = torch.from_numpy(targets_train).type(torch.LongTensor) #
# long

# PyTorch Tensor
featuresTest = torch.from_numpy(features_test) #
```


```
targetsTest = torch.from_numpy(targets_test).type(torch.LongTensor) #
↳ long
```

```
[29]: # batch_size epoch
batch_size = 100 #
n_iters = 10000 #
num_epochs = n_iters / (len(features_train) / batch_size) # epoch
num_epochs = int(num_epochs)

# PyTorch
train = TensorDataset(featuresTrain, targetsTrain) #
test = TensorDataset(featuresTest, targetsTest) #

# DataLoader
train_loader = DataLoader(train, batch_size=batch_size, shuffle=False) #
test_loader = DataLoader(test, batch_size=batch_size, shuffle=False) #
print(len(features_train))
#
plt.imshow(features_numpy[10].reshape(28, 28)) # 10 28x28
plt.axis("off") #
plt.title(str(targets_numpy[10])) #
plt.savefig('graph.png') #
plt.show() #
```

33600



```

[30]: class LSTMModel(nn.Module):
    def __init__(self, input_dim, hidden_dim, layer_dim, output_dim):
        super(LSTMModel, self).__init__()

        #
        self.hidden_dim = hidden_dim

        # LSTM
        self.layer_dim = layer_dim

        # LSTM
        self.lstm = nn.LSTM(input_dim, hidden_dim, layer_dim, batch_first=True)

        #
        self.fc = nn.Linear(hidden_dim, output_dim)

    def forward(self, x):

        #
        h0 = Variable(torch.zeros(self.layer_dim, x.size(0), self.hidden_dim))
        ↪ #
        c0 = Variable(torch.zeros(self.layer_dim, x.size(0), self.hidden_dim))
        ↪ #

        # LSTM
        out, (hn, cn) = self.lstm(x, (h0, c0))

        #
        out = self.fc(out[:, -1, :])
        return out

#
input_dim = 28 # 28
hidden_dim = 100 #
layer_dim = 1 #
output_dim = 10 # 0-9

# RNN
model = LSTMModel(input_dim, hidden_dim, layer_dim, output_dim)

#
error = nn.CrossEntropyLoss()

```

```
# SGD
learning_rate = 0.05
optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate)
```

```
[31]: seq_dim = 28 # RNN 28

#
loss_list = [] #
iteration_list = [] #
accuracy_list = [] #
count = 0 #

#
for epoch in range(num_epochs):
    for i, (images, labels) in enumerate(train_loader):

        # RNN (batch_size, seq_dim, input_dim)
        train = Variable(images.view(-1, seq_dim, input_dim))
        labels = Variable(labels)

        #
        optimizer.zero_grad()

        #
        outputs = model(train)

        #
        loss = error(outputs, labels)

        #
        loss.backward()

        #
        optimizer.step()

        count += 1

    # 250
    if count % 250 == 0:
        correct = 0
        total = 0

        #
        for images, labels in test_loader:
            images = Variable(images.view(-1, seq_dim, input_dim))

            #
```

```

        outputs = model(images)

        #
        predicted = torch.max(outputs.data, 1)[1]

        #
        total += labels.size(0)

        #
        correct += (predicted == labels).sum()

    #
    accuracy = 100 * correct / float(total)

    #
    loss_list.append(loss.item())
    iteration_list.append(count)
    accuracy_list.append(accuracy)

    # 500
    if count % 500 == 0:
        print('Iteration: {} Loss: {} Accuracy: {} %'.format(count,
↪loss.item(), accuracy))

```

```

Iteration: 500 Loss: 2.291813850402832 Accuracy: 13.476190567016602 %
Iteration: 1000 Loss: 2.2829360961914062 Accuracy: 18.714284896850586 %
Iteration: 1500 Loss: 1.9491546154022217 Accuracy: 32.345237731933594 %
Iteration: 2000 Loss: 0.9706793427467346 Accuracy: 68.60713958740234 %
Iteration: 2500 Loss: 0.635261595249176 Accuracy: 82.61904907226562 %
Iteration: 3000 Loss: 0.2595059275627136 Accuracy: 90.83333587646484 %
Iteration: 3500 Loss: 0.356059730052948 Accuracy: 90.33333587646484 %
Iteration: 4000 Loss: 0.04518391564488411 Accuracy: 93.92857360839844 %
Iteration: 4500 Loss: 0.13296236097812653 Accuracy: 94.75 %
Iteration: 5000 Loss: 0.11834557354450226 Accuracy: 95.46428680419922 %
Iteration: 5500 Loss: 0.14057557284832 Accuracy: 95.47618865966797 %
Iteration: 6000 Loss: 0.2393888682126999 Accuracy: 95.86904907226562 %
Iteration: 6500 Loss: 0.10809016972780228 Accuracy: 96.53571319580078 %
Iteration: 7000 Loss: 0.0670095682144165 Accuracy: 96.58333587646484 %
Iteration: 7500 Loss: 0.08491396903991699 Accuracy: 96.38095092773438 %
Iteration: 8000 Loss: 0.19161094725131989 Accuracy: 96.96428680419922 %
Iteration: 8500 Loss: 0.01168085914105177 Accuracy: 96.92857360839844 %
Iteration: 9000 Loss: 0.09191018342971802 Accuracy: 96.89286041259766 %
Iteration: 9500 Loss: 0.01728264056146145 Accuracy: 97.11904907226562 %

```

```

[32]: #
plt.plot(iteration_list, loss_list)
plt.xlabel("Number of iteration")

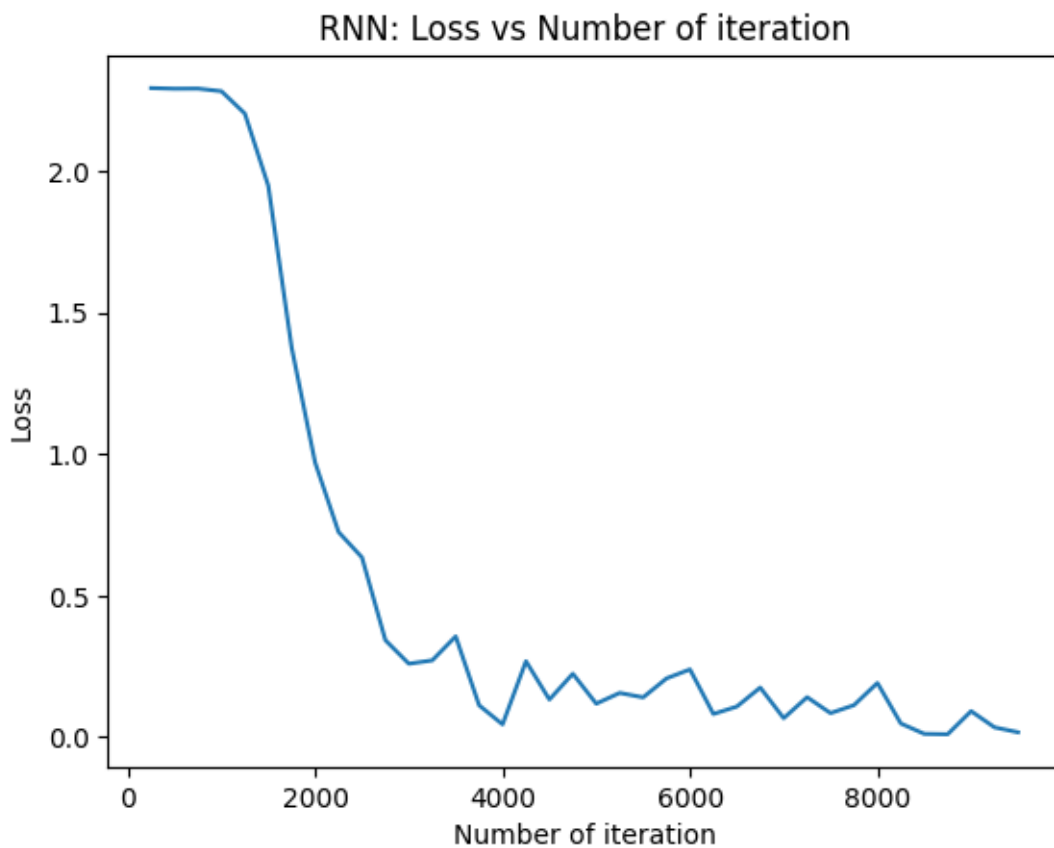
```


```

plt.ylabel("Loss")
plt.title("RNN: Loss vs Number of iteration")
plt.show()

# %%
#
plt.plot(iteration_list, accuracy_list, color="red")
plt.xlabel("Number of iteration")
plt.ylabel("Accuracy")
plt.title("RNN: Accuracy vs Number of iteration")
plt.savefig('graph.png')
plt.show()

```





```
[33]: #
model.eval()

#
with torch.no_grad():
    correct = 0
    total = 0

    # test_loader
    for images, labels in test_loader:
        images = images.view(-1, 28, 28) # (batch_size, seq_dim,
        ↪ input_dim)

        #
        outputs = model(images)

        #
        _, predicted = torch.max(outputs.data, 1)

        #
```

```

        total += labels.size(0)
        correct += (predicted == labels).sum().item() # Tensor

#
accuracy = 100 * correct / total
print(f'Test Accuracy: {accuracy} %')

```

Test Accuracy: 97.01190476190476 %

```

[36]: # test.csv
test_data = pd.read_csv(r"./input/test.csv", dtype=np.float32)

test_features_numpy = test_data.values / 255 # 0-255 -> 0-1

# test PyTorch Tensor
test_features = torch.from_numpy(test_features_numpy)

# DataLoader
test_loader = DataLoader(test_features, batch_size=9, shuffle=False)

# 2x3 6
fig, axes = plt.subplots(3, 3, figsize=(9, 9))

# 6
with torch.no_grad():
    images = next(iter(test_loader)) # DataLoader
    images = images.view(-1, 28, 28) # (batch_size, 28, 28)

    for i, ax in enumerate(axes.flat):
        ax.imshow(images[i], cmap='gray') # i
        ax.axis('off') #

plt.show() # 6

#
#
model.eval()

#
with torch.no_grad():
    images = Variable(images.view(-1, 28, 28)) # Variable (batch_size,
↪seq_dim, input_dim)


    #
    outputs = model(images)

#

```



```
_, predicted = torch.max(outputs.data, 1)
# 4.
print("Predicted Labels:", predicted.numpy()) #
```



Predicted Labels: [2 0 9 9 3 9 0 3 0]